

Job Queueing and Submission

Introduction

Habana uses Slurm for job submission and queueing. Below are some of the important commands for using Slurm. For more information refer to [Slurm Documentation](#).

Note: Run the python scripts using `srun` or `sbatch`, to ensure that concurrent jobs do not interfere with each other.

`srun`

The Slurm command `srun` can be used to run individual python scripts in parallel with other scripts on a cluster managed by Slurm. Examples of `srun` usage are shown below.

```
srun python demo_mnist.py --hpu
```

`sbatch`

Alternatively, jobs can be submitted to the Slurm workload manager through a batch script by using `sbatch` command. To do this, create a bash script (submit-demo-mnist-job.sh here as an example) with the commands that you want to execute.

```
#!/bin/sh  
  
srun python demo_mnist.py --hpu
```

Then pass the bash script as an input to the `sbatch` command as shown below.

```
sbatch submit-demo-mnist-job.sh
```

In case of the need to use multiple HPU's (2 in the example shown below), the `sbatch` command would be altered as:

```
sbatch --gres=hpu:2 <your_script.sh>
```

`squeue`

`Squeue` command provides information about jobs located in the Slurm scheduling queue.

```
queue
```

scancel

Scancel is used to signal or cancel jobs, job arrays or job steps.

```
scancel job_id
```