

# ResNet50 Keras Model Scaling

---

## TODO

For those who need to count HPUs in use on a habana node, this hack works:

```
echo $(hl-smi | grep "Mib \/" | grep -v 512 | wc -l)
```

## Login

On your development machine:

```
ssh -J CELSGCEUserID@homes.cels.anl.gov CELSGCEUserID@habana-01.ai.alcf.anl.gov  
CELS password
```

## Clone ResNet50

```
cd /path/to/location/for/repo  
git clone git@git.cels.anl.gov:ai-testbed-apps/habana/resnet50.git  
  
cd resnet50
```

## Create Venv

```
python3.8 -m venv --system-site-packages ~/venvs/habana/resnet50_TF_venv  
  
source ~/venvs/habana/resnet50_TF_venv/bin/activate
```

## Define PYTHON

Set the Python environment variable:

```
export PYTHON=$(which python)
```

## Define PYTHONPATH

Set the Python path:

```
export PYTHONPATH=/path/to/Model-References/:$(which python)
```

## Environment Variables

Set these:

```
export HABANA_LOGS=~/.habana_logs  
export MPI_ROOT=/usr/local/openmpi
```

## Change Directory

If necessary:

```
cd /path/to/ai-testbed-apps/habana/resnet50
```

## Install Requirements

```
cd /path/to/ai-testbed-apps/habana/resnet50/Resnets/resnet_keras
```

Install required packages using pip:

```
python3 -m pip install -r requirements.txt
```

## Training data

Skip this because it has already been done.

Also, the commands may not be quite right.

The ResNet50 Keras script operates on ImageNet 1k, a widely popular image classification dataset from the ILSVRC challenge. In order to obtain the dataset, follow these steps which can take over twelve hours:

1. Sign up with <http://image-net.org/download-images> and acquire the rights to download original images
2. Follow the link to the 2012 ILSVRC and download `ILSVRC2012_img_val.tar` and `ILSVRC2012_img_train.tar`.
3. Use the commands below - they will prepare the dataset under `/lambda_stor/data/imagenet/tf_records`. This is the default `data_dir` for the training script. In `/lambda_stor/data/imagenet/tf_records/train` and `/lambda_stor/data/imagenet/tf_records/val` directories original JPEG files will stay and can

be used for Media Loading Acceleration on Gaudi2. See examples with `--data_dir` and `--jpeg_data_dir` parameters for details.

```
export IMAGENET_HOME=/lambda_stor/data/imagenet/tf_records
export PYTHON=$(which python)
source ~/venvs/habana/resnet50_TF_venv/bin/activate
mkdir -p $IMAGENET_HOME/validation
mkdir -p $IMAGENET_HOME/train
tar xf ILSVRC2012_img_val.tar -C $IMAGENET_HOME/validation
tar xf ILSVRC2012_img_train.tar -C $IMAGENET_HOME/train
cd $IMAGENET_HOME/train
for f in *.tar; do
    d=`basename $f .tar`
    mkdir $d
    tar xf $f -C $d
done
cd $IMAGENET_HOME
rm $IMAGENET_HOME/train/*.tar # optional
wget -O synset_labels.txt
https://raw.githubusercontent.com/tensorflow/models/master/research/slim/datasets/imagenet_2012_validation_synset_labels.txt
cd /path/to/Model-References/TensorFlow/computer_vision/Resnets
python3 preprocess_imagenet.py \
    --raw_data_dir=$IMAGENET_HOME \
    --local_scratch_dir=$IMAGENET_HOME/tf_records
mkdir -p $IMAGENET_HOME/val
mv $IMAGENET_HOME/validation/* $IMAGENET_HOME/val
cd $IMAGENET_HOME/val
wget -qO-
https://raw.githubusercontent.com/soumith/imagenetloader.torch/master/valprep.sh | bash
```

## Single-Server Habana Resnet Training Steps

### Using LARS Optimizer

Using LARS optimizer usually requires changing the default values of some hyperparameters and should be manually set for `resnet_ctl_imagenet_main.py`. The recommended parameters together with their default values are presented below:

Parameter	Value
optimizer	LARS
base_learning_rate	2.5 or 9.5*
warmup_epochs	3
lr_schedule	polynomial
label_smoothing	0.1

Parameter	Value
weight_decay	0.0001
single_l2_loss_op	True

\*2.5 is the default value for single card (1 HPU) trainings, otherwise, the default is 9.5. These values have been determined experimentally.

One HPU on 1 server, batch 256, 5 epochs, BF16 precision, LARS

Scaling Runs

```
time $PYTHON resnet_ctl_imagenet_main.py \
  -dt bf16 \
  --data_loader_image_type bf16 \
  -te 5 \
  -bs 256 \
  --optimizer LARS \
  --base_learning_rate 9.5 \
  --warmup_epochs 3 \
  --lr_schedule polynomial \
  --label_smoothing 0.1 \
  --weight_decay 0.0001 \
  --single_l2_loss_op \
  --data_dir /lambda_stor/data/imagenet/tf_records
```

Statistics:

```
real    67m37.184s
'avg_exp_per_second': 1594.6571942068174
```

Two HPUs on 1 server, batch 256, 5 epochs, BF16 precision, LARS

```
time mpirun \
  --allow-run-as-root \
  --bind-to core \
  -np 2 \
  --map-by socket:PE=7 \
  --merge-stderr-to-stdout \
  $PYTHON resnet_ctl_imagenet_main.py \
    --dtype bf16 \
    --data_loader_image_type bf16 \
    --use_horovod \
    -te 5 \
    -bs 256 \
    --optimizer LARS \
    --base_learning_rate 9.5 \
```

```
--warmup_epochs 3 \
--lr_schedule polynomial \
--label_smoothing 0.1 \
--weight_decay 0.0001 \
--single_l2_loss_op \
--data_dir /lambda_stor/data/imagenet/tf_records
```

Statistics:

```
real      36m48.760s
'avg_exp_per_second': 2990.6616862467968
```

Four HPU's on 1 server, batch 256, 5 epochs, BF16 precision, LARS

```
time mpirun \
  --allow-run-as-root \
  --bind-to core \
  -np 4 \
  --map-by socket:PE=7 \
  --merge-stderr-to-stdout \
  $PYTHON resnet_ctl_imagenet_main.py \
    --dtype bf16 \
    --data_loader_image_type bf16 \
    --use_horovod \
    -te 5 \
    -bs 256 \
    --optimizer LARS \
    --base_learning_rate 9.5 \
    --warmup_epochs 3 \
    --lr_schedule polynomial \
    --label_smoothing 0.1 \
    --weight_decay 0.0001 \
    --single_l2_loss_op \
    --data_dir /lambda_stor/data/imagenet/tf_records
```

Statistics:

```
real      19m34.556s
'avg_exp_per_second': 5708.747706392027
```

Eight HPU's on 1 server, batch 256, 5 epochs, BF16 precision, LARS

```
time mpirun \
  --allow-run-as-root \
  --bind-to core \
```

```

-np 8 \
--map-by socket:PE=7 \
--merge-stderr-to-stdout \
  $PYTHON resnet_ctl_imagenet_main.py \
    --dtype bf16 \
    --data_loader_image_type bf16 \
    --use_horovod \
    -te 5 \
    -bs 256 \
    --optimizer LARS \
    --base_learning_rate 9.5 \
    --warmup_epochs 3 \
    --lr_schedule polynomial \
    --label_smoothing 0.1 \
    --weight_decay 0.0001 \
    --single_l2_loss_op \
    --data_dir /lambda_stor/data/imagenet/tf_records

```

Statistics:

```

real    11m11.386s
'avg_exp_per_second': 10484.725513908621

```

## Multi-Server Habana Resnet Training Steps

One-time per user ssh key set up

On both Habana machines set up the ssh key

### Habana-01

On **Habana-01**

```

mkdir ~/.ssh
cd ~/.ssh
ssh-keygen -t rsa -b 4096
#Accept default filename of id_rsa
#Enter passphrase (empty for no passphrase):
#Enter same passphrase again:
cat id_rsa.pub >> authorized_keys

```

```
ssh-keyscan -H 140.221.77.101 >> ~/.ssh/known_hosts
```

You should see:

```
# 140.221.77.101:22 SSH-2.0-OpenSSH_8.2p1 Ubuntu-4ubuntu0.5
# 140.221.77.101:22 SSH-2.0-OpenSSH_8.2p1 Ubuntu-4ubuntu0.5
# 140.221.77.101:22 SSH-2.0-OpenSSH_8.2p1 Ubuntu-4ubuntu0.5
# 140.221.77.101:22 SSH-2.0-OpenSSH_8.2p1 Ubuntu-4ubuntu0.5
# 140.221.77.101:22 SSH-2.0-OpenSSH_8.2p1 Ubuntu-4ubuntu0.5
```

```
ssh-keyscan -H 140.221.77.102 >> ~/.ssh/known_hosts
```

You should see:

```
# 140.221.77.102:22 SSH-2.0-OpenSSH_8.2p1 Ubuntu-4ubuntu0.5
# 140.221.77.102:22 SSH-2.0-OpenSSH_8.2p1 Ubuntu-4ubuntu0.5
# 140.221.77.102:22 SSH-2.0-OpenSSH_8.2p1 Ubuntu-4ubuntu0.5
# 140.221.77.102:22 SSH-2.0-OpenSSH_8.2p1 Ubuntu-4ubuntu0.5
# 140.221.77.102:22 SSH-2.0-OpenSSH_8.2p1 Ubuntu-4ubuntu0.5
```

Verify you can ssh from habana-01 to habana-02 without a password.

Verify vice-versa.

## ResNet50

Ensure \$MPI\_ROOT is set

```
echo $MPI_ROOT
```

If not,

```
export MPI_ROOT=/usr/local/openmpi
```

## Tensorflow

Test with:

```
cd /path/to/Model-References/TensorFlow/computer_vision/Resnets/resnet_keras
```

Verify a connection with

```
mpirun -v --allow-run-as-root --mca btl_tcp_if_include enp75s0f0,ens1f0 --
tag-output --merge-stderr-to-stdout --prefix /usr/local/openmpi --host
140.221.77.101:8,140.221.77.102:8 hostname
```

Almost immediately, you should see

```
[1,0]<stdout>:habana-01
[1,1]<stdout>:habana-01
[1,2]<stdout>:habana-01
[1,3]<stdout>:habana-01
[1,4]<stdout>:habana-01
[1,5]<stdout>:habana-01
[1,6]<stdout>:habana-01
[1,7]<stdout>:habana-01
[1,8]<stdout>:habana-02
[1,9]<stdout>:habana-02
[1,10]<stdout>:habana-02
[1,11]<stdout>:habana-02
[1,12]<stdout>:habana-02
[1,13]<stdout>:habana-02
[1,14]<stdout>:habana-02
[1,15]<stdout>:habana-02
```

If the above command hangs, there is a communication problem.

Sixteen HPUs on 2 servers, batch 256, 5 epochs, BF16 precision, LARS

```
mpirun -v \
--allow-run-as-root \
-np 16 \
--mca btl_tcp_if_include 192.168.201.0/24 \
--tag-output \
--merge-stderr-to-stdout \
--prefix /usr/local/openmpi \
-H 140.221.77.101:8,140.221.77.102:8 \
-x GC_KERNEL_PATH \
-x HABANA_LOGS \
-x PYTHONPATH \
-x HCCL_SOCKET_IFNAME=enp75s0f0,ens1f0 \
-x HCCL_OVER_TCP=1 \
$PYTHON resnet_ctl_imagenet_main.py \
-dt bf16 \
-dlit bf16 \
-bs 256 \
-te 5 \
--use_horovod \
--data_dir /lambda_stor/data/imagenet/tf_records/ \
--optimizer LARS \
```



```
--base_learning_rate 9.5 \
--warmup_epochs 3 \
--lr_schedule polynomial \
--label_smoothing 0.1 \
--weight_decay 0.0001 \
--single_l2_loss_op \
--model_dir=`mktemp -d`
```

#### Statistics:

```
2022-09-27 18:23:10.461103
2022-09-27 18:31:34.638972
real      8:24
'avg_exp_per_second': 15481.107839604083
```

### PyTorch

This section has not been tested.

```
` ``bash
cd Model-References/PyTorch/computer_vision/classification/torchvision/
export MASTER_PORT=12355
export MASTER_ADDR=192.168.201.101
mpirun --allow-run-as-root --mca --bind-to core --map-by ppr:4:socket:PE=7
-np 16 --mca btl_tcp_if_include 192.168.201.0/24 --merge-stderr-to-stdout -
-prefix $MPI_ROOT -H 140.221.77.101:8,140.221.77.102:8 -x GC_KERNEL_PATH -x
PYTHONPATH -x MASTER_ADDR -x MASTER_PORT -x
HCCL_SOCKET_IFNAME=enp75s0f0,ens1f0 -x HCCL_OVER_TCP=1 $PYTHON -u train.py
--batch-size=256 --model=resnet50 --device=hpu --workers=8 --print-freq=1 -
-deterministic --data-path=/lambda_stor/habana/data/tensorflow/imagenet --
epochs=2 --hmp --hmp-bf16 ./ops_bf16_Resnet.txt --hmp-fp32
./ops_fp32_Resnet.txt --custom-lr-values 0.475
```